

Little Book of Semaphores, Chapter 3

Geoffrey Matthews
Western Washington University

January 20, 2016

Signaling

Initialization

```
1 sem = Semaphore(0)
```


Thread A

```
1 a1  
2 sem.signal()
```

Thread B


```
1 sem.wait()  
2 b1
```

Rendezvous



- A has to wait for b1 to finish before a2
- B has to wait for a1 to finish before b2
- Ideas?

Rendezvous



- A has to wait for b1 to finish before a2
- B has to wait for a1 to finish before b2
- Ideas?
- Hint: use `aArrived` and `bArrived`

Rendezvous

Initialization

```
1  aArrived = Semaphore(0)
2  bArrived = Semaphore(0)
```

Thread A

```
1  a1
2  aArrived.signal()
3  bArrived.wait()
4  a2
```

Thread B

```
1  b1
2  bArrived.signal()
3  aArrived.wait()
4  b2
```

Rendezvous 2

Thread A

```
1 a1
2 bArrived.wait()
3 aArrived.signal()
4 a2
```

Thread B

```
1 b1
2 bArrived.signal()
3 aArrived.wait()
4 b2
```

- Will it still work?

Rendezvous 2

Thread A

```
1 a1
2 bArrived.wait()
3 aArrived.signal()
4 a2
```

Thread B

```
1 b1
2 bArrived.signal()
3 aArrived.wait()
4 b2
```

- Will it still work?
- Try to limit context switches.
- Do everything you can before waiting.

Rendezvous 3

Thread A

```
1 a1
2 bArrived.wait()
3 aArrived.signal()
4 a2
```

Thread B

```
1 b1
2 aArrived.wait()
3 bArrived.signal()
4 b2
```

- Will it still work?

Rendezvous 3

Thread A

```
1 a1
2 bArrived.wait()
3 aArrived.signal()
4 a2
```

Thread B

```
1 b1
2 aArrived.wait()
3 bArrived.signal()
4 b2
```

- Will it still work?
- **Deadlock!**

Mutex




- Any problems?

Mutex




- Any problems?
- Cannot assume operations are atomic.
- Even if we write:
count++

Mutex



- Can't let both threads operate at the same time.
- Don't care which goes first.
- Ideas?

Mutex



- Can't let both threads operate at the same time.
- Don't care which goes first.
- Ideas?
- Hint: create a semaphore `mutex` initialized to 1.
- The rock in the box!

Mutex

Initialization

```
1 mutex = Semaphore(1)
```

Thread A

```
1 mutex.wait()  
2 # critical section  
3 count = count + 1  
4 mutex.signal()
```

Thread B

```
1 mutex.wait()  
2 # critical section  
3 count = count + 1  
4 mutex.signal()
```

Mutex

Initialization

```
1 mutex = Semaphore(1)
```

Thread A

```
1 mutex.wait()  
2 # critical section  
3 count = count + 1  
4 mutex.signal()
```

Thread B

```
1 mutex.wait()  
2 # critical section  
3 count = count + 1  
4 mutex.signal()
```

- A **symmetric** solution.
- Symmetric solutions are easy to generalize.

Mutex

Initialization

```
1 mutex = Semaphore(1)
```

Thread A

```
1 mutex.wait()  
2 # critical section  
3 count = count + 1  
4 mutex.signal()
```

Thread B

```
1 mutex.wait()  
2 # critical section  
3 count = count + 1  
4 mutex.signal()
```

- A **symmetric** solution.
- Symmetric solutions are easy to generalize.
- Metaphorically we can look at this as a **token**:
the rock in the box, the talking stick

Mutex

Initialization

```
1 mutex = Semaphore(1)
```

Thread A

```
1 mutex.wait()  
2 # critical section  
3 count = count + 1  
4 mutex.signal()
```

Thread B

```
1 mutex.wait()  
2 # critical section  
3 count = count + 1  
4 mutex.signal()
```

- A **symmetric** solution.
- Symmetric solutions are easy to generalize.
- Metaphorically we can look at this as a **token**:
the rock in the box, the talking stick
- Another metaphor is a **lock**
- Sometimes called “getting” and “releasing” a lock.

Multiplex

- Generalize the mutex so that at most n threads can access the critical section at a time.
- Ideas?

Multiplex

- Generalize the mutex so that at most n threads can access the critical section at a time.
- Ideas?
- Initialize the mutex to n .

Initialization

```
1 mutex = Semaphore(n)
```

Thread i

```
1 mutex.wait()  
2 # critical section  
3 count = count + 1  
4 mutex.signal()
```

Thread j

```
1 mutex.wait()  
2 # critical section  
3 count = count + 1  
4 mutex.signal()
```

Barrier

Initialization

```
1  aArrived = Semaphore(0)
2  bArrived = Semaphore(0)
```

Thread A


```
1  a1
2  aArrived.signal()
3  bArrived.wait()
4  a2
```

Thread B

```
1  b1
2  bArrived.signal()
3  aArrived.wait()
4  b2
```

- Recall the rendezvous, above.
- Is there a way to generalize this to n threads?

Barrier



- We want all tasks to finish rendezvous before beginning critical point
- When the first $n - 1$ threads arrive, they should block until the n th thread arrives, when all should proceed.
- Ideas?

Barrier Hint

Initialization

```
1 n = number of threads
2 count = 0
3 mutex = Semaphore(1)
4 barrier = Semaphore(0)
```

- count keeps track of how many threads have arrived
- mutex provides atomic increment of count
- barrier is locked until all threads arrive

Barrier non-solution

Thread i

```
1 #rendezvous
2 mutex.wait()
3   count = count + 1
4 mutex.signal()
5
6 if count == n:    barrier.signal()
7
8 barrier.wait()
9
10 #critical point
```

- What's wrong?

Barrier non-solution

Thread i

```
1 #rendezvous
2 mutex.wait()
3   count = count + 1
4 mutex.signal()
5
6 if count == n:    barrier.signal()
7
8 barrier.wait()
9
10 #critical point
```

- What's wrong?
- Deadlock!

Barrier non-solution

Thread i

```
1 #rendezvous
2 mutex.wait()
3   count = count + 1
4 mutex.signal()
5
6 if count == n:    barrier.signal()
7
8 barrier.wait()
9
10 #critical point
```

- What's wrong?
- Deadlock!
- Does it *always* deadlock?

Barrier working solution

Thread i

```
1 #rendezvous
2 mutex.wait()
3     count = count + 1
4 mutex.signal()
5
6 if count == n: barrier.signal()
7
8 barrier.wait()
9 barrier.signal()
10
11 #critical point
```

- wait then signal is called a **turnstile**
- It blocks all threads until one thread signals, then lets all through.
- An “open” turnstile can be *locked* by calling wait and *unlocked* by calling signal.

Barrier working solution

Thread i

```
1 #rendezvous
2 mutex.wait()
3   count = count + 1
4 mutex.signal()
5
6 if count == n: barrier.signal()
7
8 barrier.wait()
9 barrier.signal()
10
11 #critical point
```

- After the n th thread, what state is the turnstile in?
- Is the barrier reusable? Can we put this code in a loop?

Another barrier non-solution

Thread i

```
1 #rendezvous
2
3 mutex.wait()
4     count = count + 1
5     if count == n: barrier.signal()
6
7     barrier.wait()
8     barrier.signal()
9 mutex.signal()
10
11 #critical point
```

- Deadlock again. Why?

Another barrier non-solution

Thread i

```
1 #rendezvous
2
3 mutex.wait()
4     count = count + 1
5     if count == n: barrier.signal()
6
7     barrier.wait()
8     barrier.signal()
9 mutex.signal()
10
11 #critical point
```

- Deadlock again. Why?
- Common source of deadlocks: blocking on a semaphore while holding a mutex.

Reusable barrier non-solution #1

Thread i

```
1 #rendezvous
2 mutex.wait(); count += 1; mutex.signal()
3 if count == n: turnstile.signal()
4 turnstile.wait()
5 turnstile.signal()
6 #critical point
7 mutex.wait(); count -= 1; mutex.signal()
8 if count == 0: turnstile.wait()
```

- What's wrong?

Reusable barrier non-solution #1

```
Thread i
1 #rendezvous
2 mutex.wait(); count += 1; mutex.signal()
3 if count == n: turnstile.signal()
4 turnstile.wait()
5 turnstile.signal()
6 #critical point
7 mutex.wait(); count -= 1; mutex.signal()
8 if count == 0: turnstile.wait()
```

- What's wrong?
- If we interrupt a process *just before* evaluating the first conditional?

Reusable barrier non-solution #1

```
Thread i
1 #rendezvous
2 mutex.wait(); count += 1; mutex.signal()
3 if count == n: turnstile.signal()
4 turnstile.wait()
5 turnstile.signal()
6 #critical point
7 mutex.wait(); count -= 1; mutex.signal()
8 if count == 0: turnstile.wait()
```

- What's wrong?
- If we interrupt a process *just before* evaluating the first conditional?
- It is possible that *all* the threads will see `count == n` and signal the turnstile.

Reusable barrier non-solution #1

```
Thread i
1 #rendezvous
2 mutex.wait(); count += 1; mutex.signal()
3 if count == n: turnstile.signal()
4 turnstile.wait()
5 turnstile.signal()
6 #critical point
7 mutex.wait(); count -= 1; mutex.signal()
8 if count == 0: turnstile.wait()
```

- What's wrong?
- If we interrupt a process *just before* evaluating the first conditional?
- It is possible that *all* the threads will see `count == n` and signal the turnstile.
- If the second conditional is interrupted?

Reusable barrier non-solution #1

```
Thread i
1 #rendezvous
2 mutex.wait(); count += 1; mutex.signal()
3 if count == n: turnstile.signal()
4 turnstile.wait()
5 turnstile.signal()
6 #critical point
7 mutex.wait(); count -= 1; mutex.signal()
8 if count == 0: turnstile.wait()
```

- What's wrong?
- If we interrupt a process *just before* evaluating the first conditional?
- It is possible that *all* the threads will see `count == n` and signal the turnstile.
- If the second conditional is interrupted?
- Deadlock!

Reusable barrier non-solution #2

Thread i

```
1 rendezvous
2 mutex.wait();
3   count += 1;
4   if count == n: turnstile.signal()
5 mutex.signal()
6 turnstile.wait()
7 turnstile.signal()
8 #critical point
9 mutex.wait();
10   count -= 1;
11   if count == 0: turnstile.wait()
12 mutex.signal()
```

- Still doesn't work. Why?

Reusable barrier non-solution #2

Thread i

```
1 #rendezvous
2 mutex.wait();
3   count += 1;
4   if count == n: turnstile.signal()
5 mutex.signal()
6 turnstile.wait()
7 turnstile.signal()
8 #critical point
9 mutex.wait();
10   count -= 1;
11   if count == 0: turnstile.wait()
12 mutex.signal()
```

- Still doesn't work. Why?
- Hint: this is meant to be in a loop.

Reusable barrier non-solution #2

Thread i

```
1 #rendezvous
2 mutex.wait();
3   count += 1;
4   if count == n: turnstile.signal()
5 mutex.signal()
6 turnstile.wait()
7 turnstile.signal()
8 #critical point
9 mutex.wait();
10   count -= 1;
11   if count == 0: turnstile.wait()
12 mutex.signal()
```

- Still doesn't work. Why?
- Hint: this is meant to be in a loop.
- One thread could go around and through the turnstile again while the others sit there, getting one lap ahead.
- What to do?

Reusable barrier solution

Thread i

```
1 #rendezvous
2 mutex.wait();
3     count += 1;
4     if count == n:
5         turnstile2.wait()           # lock the second
6         turnstile.signal()          # unlock the first
7 mutex.signal()
8 turnstile.wait()                    # first turnstile
9 turnstile.signal()
10 #critical point
11 mutex.wait();
12     count -= 1;
13     if count == 0:
14         turnstile.wait()           # lock the first
15         turnstile2.signal()        # unlock the second
16 mutex.signal()
17 turnstile2.wait()                  # second turnstile
18 turnstile2.signal()
```

Reusable barrier solution

- Called a **two-phase barrier**
- Forces all threads to wait twice: once for all to arrive, and again for all threads to execute the critical section.

Reusable barrier solution

- Called a **two-phase barrier**
- Forces all threads to wait twice: once for all to arrive, and again for all threads to execute the critical section.
- Typical of semaphores—complex and difficult to understand.
- Can we prove it correct?

Reusable barrier solution

- Called a **two-phase barrier**
- Forces all threads to wait twice: once for all to arrive, and again for all threads to execute the critical section.
- Typical of semaphores—complex and difficult to understand.
- Can we prove it correct?
- Only the n th thread can unlock turnstiles.
- Before a thread can unlock the first turnstile, it has to close the second, and vice versa. It is therefore impossible for one thread to get ahead of the others by more than one turnstile.

Preloaded turnstile

```
1 # rendezvous
2 mutex.wait()
3   count  += 1
4   if count == n:
5       turnstile2.wait()    # lock the second
6       turnstile.signal(n)  # unlock the first
7 mutex.signal()
8 turnstile.wait()           # first turnstile
9 # critical point
10 mutex.wait()
11   count -= 1
12   if count == 0:
13       turnstile.wait()     # lock the first
14       turnstile2.signal(n) # unlock the second
15 mutex.signal()
16 turnstile2.wait()          # second turnstile
```

- Signals n at a time
- Could be done in a loop?

Barrier Object

```
1 class Barrier:
2     def __init__(self, n):
3         self.n = n
4         self.count = 0
5         self.mutex = Semaphore(1)
6         self.turnstile = Semaphore(0)
7         self.turnstile2 = Semaphore(0)
8     def phase1(self):
9         self.mutex.wait()
10        self.count += 1
11        if self.count == self.n:
12            self.turnstile2.wait()
13            self.turnstile.signal(self.n)
14        self.mutex.signal()
15        self.turnstile.wait()
16    def phase2(self):
17        self.mutex.wait()
18        self.count -= 1
19        if self.count == 0:
20            self.turnstile.wait()
21            self.turnstile2.signal(self.n)
22        self.mutex.signal()
23        self.turnstile2.wait()
```

```
1 barrier = Barrier(n)
2 ...
3 # rendezvous
4 barrier.phase1()
5 # critical point
6 barrier.phase2()
```

```
1 barrier = Barrier(n)
2 ...
3 # rendezvous
4 barrier.phase1()
5 barrier.phase2()
6 # critical point
```

Multiple barriers

```
1 class Barrier:
2     def __init__(self, n):
3         self.n = n
4         self.count = 0
5         self.mutex = Semaphore(1)
6         self.turnstile = Semaphore(0)
7         self.turnstile2 = Semaphore(0)
8     def phase1(self):
9         self.mutex.wait()
10        self.count += 1
11        if self.count == self.n:
12            self.turnstile2.wait()
13            self.turnstile.signal(self.n)
14        self.mutex.signal()
15        self.turnstile.wait()
16    def phase2(self):
17        self.mutex.wait()
18        self.count -= 1
19        if self.count == 0:
20            self.turnstile.wait()
21            self.turnstile2.signal(self.n)
22        self.mutex.signal()
23        self.turnstile2.wait()
```

```
1 barrier = Barrier(n)
2 ...
3 loop:
4     # section A
5     barrier.phase1()
6     # section B
7     barrier.phase2()
8     # section C
9     barrier.phase1()
10    # section D
11    barrier.phase2()
```

Multiple barriers

```
1 class Barrier:
2     def __init__(self, n):
3         self.n = n
4         self.count = 0
5         self.mutex = Semaphore(1)
6         self.turnstile = Semaphore(0)
7         self.turnstile2 = Semaphore(0)
8     def phase1(self):
9         self.mutex.wait()
10        self.count += 1
11        if self.count == self.n:
12            self.turnstile2.wait()
13            self.turnstile.signal(self.n)
14        self.mutex.signal()
15        self.turnstile.wait()
16    def phase2(self):
17        self.mutex.wait()
18        self.count -= 1
19        if self.count == 0:
20            self.turnstile.wait()
21            self.turnstile2.signal(self.n)
22        self.mutex.signal()
23        self.turnstile2.wait()
```

```
1 barrier = Barrier(n)
2 ...
3 loop:
4     # section A
5     barrier.phase1()
6     # section B
7     barrier.phase2()
8     # section C
9     barrier.phase1()
10    # section D
11    barrier.phase2()
```

- What if we have an odd number of sections?

Queue

- Ballroom dancing.
- If a leader arrives, it checks for a follower, if none available waits, otherwise proceeds.
- If a follower arrives, it checks for a leader, if none available waits, otherwise proceeds.
- Ideas?

Queue

- Ballroom dancing.
- If a leader arrives, it checks for a follower, if none available waits, otherwise proceeds.
- If a follower arrives, it checks for a leader, if none available waits, otherwise proceeds.
- Ideas?
- Hint:

```
1 leaderQueue = Semaphore(0)
2 followerQueue = Semaphore(0)
```

Queue solution

Leader

```
1 followerQueue.signal()  
2 leaderQueue.wait()  
3 dance()
```

Follower

```
1 leaderQueue.signal()  
2 followerQueue.wait()  
3 dance()
```


Queue solution

Leader

```
1 followerQueue.signal()  
2 leaderQueue.wait()  
3 dance()
```

Follower

```
1 leaderQueue.signal()  
2 followerQueue.wait()  
3 dance()
```

- Do the leaders and followers proceed together?

Queue solution

Leader

```
1 followerQueue.signal()  
2 leaderQueue.wait()  
3 dance()
```

Follower

```
1 leaderQueue.signal()  
2 followerQueue.wait()  
3 dance()
```

- Do the leaders and followers proceed together?
- It is possible for 100 leaders to dance before any followers do.

Queue solution

Leader

```
1 followerQueue.signal()  
2 leaderQueue.wait()  
3 dance()
```

Follower

```
1 leaderQueue.signal()  
2 followerQueue.wait()  
3 dance()
```

- Do the leaders and followers proceed together?
- It is possible for 100 leaders to dance before any followers do.
- Add constraint that only one leader and one follower can dance concurrently.
- Ideas?

Exclusive Queue Hint

```
1 leaders = followers = 0
2 mutex = Semaphore(1)
3 leaderQueue = Semaphore(0)
4 followerQueue = Semaphore(0)
5 rendezvous = Semaphore(0)
```

Exclusive Queue Solution

Leader

```
1 mutex.wait()
2 if followers > 0:
3     followers--
4     followerQueue.signal()
5 else:
6     leaders++
7     mutex.signal()
8     leaderQueue.wait()
9 dance()
10 rendezvous.wait()
11 mutex.signal()
```

Follower

```
1 mutex.wait()
2 if leaders > 0:
3     leaders--
4     leaderQueue.signal()
5 else:
6     followers++
7     mutex.signal()
8     followerQueue.wait()
9 dance()
10 rendezvous.signal()
11
```

- “Wait” means “wait on this queue”
- “Signal” means “let someone go from this queue”

FIFO Queue

Initialization

```
1 local mySem = Semaphore(0)
```

```
1 class Fifo:
2     def __init__(self):
3         self.queue = Queue()
4         self.mutex = Semaphore(1)
5
6     def wait(self, mySem):
7         self.mutex.wait()
8         self.queue.add(mySem)
9         self.mutex.signal()
10        mySem.wait()
11
12    def signal(self):
13        self.mutex.wait()
14        sem = self.queue.remove()
15        self.mutex.signal()
16        sem.signal()
```