

Little Book of Semaphores, Chapter 4

Geoffrey Matthews
Western Washington University

January 11, 2016

Producers and Consumers

Producer i

```
1 event[i] = waitForEvent()  
2 buffer.add(event[i])
```

Consumer j

```
1 event[j] = buffer.get()  
2 event[j].process()
```

- Threads must have exclusive access to the buffer. No adding and getting at the same time.
- If a consumer thread arrives when the buffer is empty, it blocks until a producer adds an item.
- `waitForEvent` and `process` can happen simultaneously, but not buffer access.
- `event` is a local variable for each thread—not shared.
- We could use an array, as above, with all producers and consumers having different indices.

Producers and Consumers Hint

```
1 mutex = Semaphore(1)
2 items = Semaphore(0)
3 local event
```

- Local events can be handled several ways:
 - Each thread has its own run-time stack. (We use this in scheme and python, where threads are functions.)
 - Threads could be objects, with local private variables.
 - Threads can use unique IDs as indices into an array.

Producer-consumer solution

Producer

```
1 event = waitForEvent()  
2 mutex.wait()  
3   buffer.add(event)  
4   items.signal()  
5 mutex.signal()
```

Consumer

```
1 items.wait()  
2 mutex.wait()  
3   event = buffer.get()  
4   mutex.signal()  
5   event.process()
```

Producer-consumer solution

Producer

```
1 event = waitForEvent()  
2 mutex.wait()  
3   buffer.add(event)  
4   items.signal()  
5 mutex.signal()
```

Consumer

```
1 items.wait()  
2 mutex.wait()  
3   event = buffer.get()  
4   mutex.signal()  
5   event.process()
```

- Could the `items.signal()` be taken out of the mutex?
- What would be the advantage?

Producer-consumer solution (slight improvement)

Producer

```
1 event = waitForEvent()  
2 mutex.wait()  
3     buffer.add(event)  
4 mutex.signal()  
5 items.signal()
```

Consumer

```
1 items.wait()  
2 mutex.wait()  
3     event = buffer.get()  
4 mutex.signal()  
5 event.process()
```

Producer-consumer solution (slight improvement)

Producer

```
1 event = waitForEvent()  
2 mutex.wait()  
3     buffer.add(event)  
4 mutex.signal()  
5 items.signal()
```

Consumer

```
1 items.wait()  
2 mutex.wait()  
3     event = buffer.get()  
4 mutex.signal()  
5 event.process()
```

- items could at times not accurately reflect the actual number of waiting consumers.

Producer-consumer solution (broken)

Producer

```
1 event = waitForEvent()  
2 mutex.wait()  
3   buffer.add(event)  
4 mutex.signal()  
5 items.signal()
```

Consumer

```
1 mutex.wait()  
2   items.wait()  
3   event = buffer.get()  
4 mutex.signal()  
5 event.process()
```

- Why is this broken?

Producer-consumer solution (broken)

Producer

```
1 event = waitForEvent()  
2 mutex.wait()  
3   buffer.add(event)  
4 mutex.signal()  
5 items.signal()
```

Consumer

```
1 mutex.wait()  
2   items.wait()  
3   event = buffer.get()  
4 mutex.signal()  
5 event.process()
```

- Why is this broken?
- Don't wait for a semaphore after grabbing a mutex!

Producer-consumer with finite buffer

Broken finite buffer solution

```
1 if items >= bufferSize:  
2     block()
```

- `items` is a semaphore, we can't check its size
- Even if we could, we could be interrupted between checking and blocking.
- We *don't* want to block inside a mutex!

Producer-consumer with finite buffer

Broken finite buffer solution

```
1 if items >= bufferSize:  
2     block()
```

- `items` is a semaphore, we can't check its size
- Even if we could, we could be interrupted between checking and blocking.
- We *don't* want to block inside a mutex!
- What to do?

Finite buffer producer-consumer hint

```
1 mutex = Semaphore(1)
2 items = Semaphore(0)
3 spaces = Semaphore(bufferSize)
```

- $\text{items} + \text{spaces} = \text{bufferSize}$

Finite buffer producer-consumer solution

```
1 mutex = Semaphore(1)
2 items = Semaphore(0)
3 spaces = Semaphore(bufferSize)
```

Producer

```
1 event = waitForEvent()
2
3 spaces.wait()
4 mutex.wait()
5     buffer.add(event)
6 mutex.signal()
7 items.signal()
```

Consumer

```
1 items.wait()
2 mutex.wait()
3     event = buffer.get()
4 mutex.signal()
5 spaces.signal()
6
7 event.process()
```

Finite buffer producer-consumer solution

```
1 mutex = Semaphore(1)
2 items = Semaphore(0)
3 spaces = Semaphore(bufferSize)
```

Producer

```
1 event = waitForEvent()
2
3 spaces.wait()
4 mutex.wait()
5     buffer.add(event)
6 mutex.signal()
7 items.signal()
```

Consumer

```
1 items.wait()
2 mutex.wait()
3     event = buffer.get()
4 mutex.signal()
5 spaces.signal()
6
7 event.process()
```

- Note that $\text{items} + \text{spaces}$ is a constant.
- This is an *invariant*
- Using invariants is a good way to design programs and help prove properties.

Readers-writers problem

- Suppose a number of processes all access the same data.
- Any number of readers can be in the critical section simultaneously.
- Writers must have exclusive access to the critical section.
- Ideas?

Readers-writers hint

```
1 readers = 0
2 mutex = Semaphore(1)
3 roomEmpty = Semaphore(1)
```

- “wait” means “wait for the condition to be true”
- “signal” means “signal that the condition is true”

Readers-writers solution

Writers

```
1 roomEmpty.wait()
2 # critical section for writer
3 roomEmpty.signal()
```

Readers

```
1 mutex.wait()
2   readers += 1
3   if readers == 1:
4       # first in locks:
5       roomEmpty.wait()
6   mutex.signal()
7
8   # critical section for reader
9
10  mutex.wait()
11   readers -= 1
12   if readers == 0:
13       # last out unlocks
14       roomEmpty.signal()
15  mutex.signal()
```

- A Lightswitch

Claims useful in a proof of correctness

- Only one reader can queue on `roomEmpty`
- Several writers might be queued on `roomEmpty`
- When a reader signals `roomEmpty` the room is empty

A lightswitch object

```
1 class Lightswitch:
2     def __init__(self):
3         self.counter = 0
4         self.mutex = Semaphore(1)
5
6     def lock(self, semaphore):
7         self.mutex.wait()
8         self.counter += 1
9         if self.counter == 1:
10             semaphore.wait()
11         self.mutex.signal()
12
13     def unlock(self, semaphore):
14         self.mutex.wait()
15         self.counter -= 1
16         if self.counter == 0:
17             semaphore.signal()
18         self.mutex.signal()
```

Initialization

```
readswitch = Lightswitch()
roomEmpty = Semaphore(1)
```

Readers

```
readswitch.lock(roomEmpty)
# critical section
readswitch.unlock(roomEmpty)
```

Writers

```
roomEmpty.wait()
# critical section for writer
roomEmpty.signal()
```

Starvation

- No deadlock in the above readers-writers solution.
- However, it is possible for a writer to **starve**.
- While a writer is blocked, readers can come and go, and the writer never progresses.
- (In the buffer problem, readers eventually empty the buffer, but we can imagine readers who simply examine the buffer without removing an item.)

Starvation

- No deadlock in the above readers-writers solution.
- However, it is possible for a writer to **starve**.
- While a writer is blocked, readers can come and go, and the writer never progresses.
- (In the buffer problem, readers eventually empty the buffer, but we can imagine readers who simply examine the buffer without removing an item.)
- Puzzle: extend the solution so that when a writer arrives, the existing readers can finish, but no additional readers may enter.

Starvation

- No deadlock in the above readers-writers solution.
- However, it is possible for a writer to **starve**.
- While a writer is blocked, readers can come and go, and the writer never progresses.
- (In the buffer problem, readers eventually empty the buffer, but we can imagine readers who simply examine the buffer without removing an item.)
- Puzzle: extend the solution so that when a writer arrives, the existing readers can finish, but no additional readers may enter.
- Hint: Add a turnstyle and allow the writers to lock it.

No-starve readers-writers hint

```
1 readSwitch = Lightswitch()  
2 roomEmpty = Semaphore(1)  
3 turnstile = Semaphore(1)
```

No-starve readers-writers hint

```
1 readSwitch = Lightswitch()  
2 roomEmpty = Semaphore(1)  
3 turnstile = Semaphore(1)
```

- turnstile is a turnstile for readers and a mutex for writers

No-starve readers-writers solution

Writers

```
1 turnstile.wait()
2   roomEmpty.wait()
3   # critical section
4 turnstile.signal()
5
6 roomEmpty.signal()
```

Readers

```
1 turnstile.wait()
2 turnstile.signal()
3
4 readSwitch.lock(roomEmpty)
5   # critical section
6 readSwitch.unlock(roomEmpty)
```

- turnstile is a turnstile for readers and a mutex for writers

No-starve readers-writers solution

Writers

```
1 turnstile.wait()
2   roomEmpty.wait()
3   # critical section
4   turnstile.signal()
5
6   roomEmpty.signal()
```

Readers

```
1 turnstile.wait()
2   turnstile.signal()
3
4   readSwitch.lock(roomEmpty)
5   # critical section
6   readSwitch.unlock(roomEmpty)
```

- turnstile is a turnstile for readers and a mutex for writers
- It is now possible for *readers* to starve!

Priority Scheduling

- Some schedulers allow priority scheduling.
- Puzzle: Write a solution to readers-writers that gives priority to writers. In other words, once a writer arrives, no readers are allowed in the critical section until *all* writers have left the system.

Priority Scheduling

- Some schedulers allow priority scheduling.
- Puzzle: Write a solution to readers-writers that gives priority to writers. In other words, once a writer arrives, no readers are allowed in the critical section until *all* writers have left the system.
- Hint: use two lightswitches

Writer-priority readers-writers hint

```
1 readSwitch = Lightswitch()  
2 writeSwitch = Lightswitch()  
3 mutex = Semaphore(1)  
4 noReaders = Semaphore(1)  
5 noWriters = Semaphore(1)
```

Writer-priority readers-writers solution

Writers

```
1 writeSwitch.lock(noReaders)
2   noWriters.wait()
3
4   # critical section
5
6   noWriters.signal()
7 writeSwitch.unlock(noReaders)
```

Readers

```
1 noReaders.wait()
2   readSwitch.lock(noWriters)
3   noReaders.signal()
4
5   # critical section
6
7   readSwitch.unlock(noWriters)
```

- Writers in critical section hold *both* noReaders and noWriters.
- writeSwitch allows writers to queue on noWriters, but keeps noReaders locked
- The last writer signals noReaders

Writer-priority readers-writers solution

Writers

```
1 writeSwitch.lock(noReaders)
2   noWriters.wait()
3
4   # critical section
5
6   noWriters.signal()
7 writeSwitch.unlock(noReaders)
```

Readers

```
1 noReaders.wait()
2   readSwitch.lock(noWriters)
3   noReaders.signal()
4
5   # critical section
6
7   readSwitch.unlock(noWriters)
```

- Writers in critical section hold *both* noReaders and noWriters.
- writeSwitch allows writers to queue on noWriters, but keeps noReaders locked
- The last writer signals noReaders
- Readers in critical section hold noWriters but don't hold noReaders, so a writer can lock noReaders
- The last reader signals noWriters so writers can go

Thread starvation

- We just addressed **categorical starvation**: one category of threads makes another category starve.
- **Thread starvation** is the more general possibility of a thread waiting indefinitely while other threads proceed.

Thread starvation

- We just addressed **categorical starvation**: one category of threads makes another category starve.
- **Thread starvation** is the more general possibility of a thread waiting indefinitely while other threads proceed.
- Part of the problem is the responsibility of the scheduler. If a thread is never scheduled, it is starved.

Thread starvation

- We just addressed **categorical starvation**: one category of threads makes another category starve.
- **Thread starvation** is the more general possibility of a thread waiting indefinitely while other threads proceed.
- Part of the problem is the responsibility of the scheduler. If a thread is never scheduled, it is starved.
- Some schedulers use algorithms that guarantee bounded waiting.

Thread starvation

- If we don't want to assume too much about the scheduler, can we assume:
- **Property 1:** if there is only one thread that is ready to run, the scheduler has to let it run.
- This would be sufficient for the boundary problem.
- In general we need a stronger assumption.

Thread starvation

- **Property 2:** if a thread is ready to run, then the time it waits until it runs is bounded.
- We use this assumption in all our work.
- Some schedulers in the real world do not guarantee this strictly.
- Property 2 is not strong enough if we use semaphores. Why?

Semaphore starvation

- The weakest assumption about semaphores that makes it possible to avoid starvation is:
- **Property 3:** if there are threads waiting on a semaphore when a thread executes `signal`, then one of the waiting threads has to be woken.

Semaphore starvation

- The weakest assumption about semaphores that makes it possible to avoid starvation is:
- **Property 3:** if there are threads waiting on a semaphore when a thread executes `signal`, then one of the waiting threads has to be woken.
- Prevents a thread from signalling a semaphore, racing around a loop and catching its own signal!

```
_____ Thread i _____  
1   while True:  
2       mutex.wait()  
3       # critical section  
4       mutex.signal()
```

Semaphore starvation

- The weakest assumption about semaphores that makes it possible to avoid starvation is:
- **Property 3:** if there are threads waiting on a semaphore when a thread executes `signal`, then one of the waiting threads has to be woken.

- Prevents a thread from signalling a semaphore, racing around a loop and catching its own signal!

```
1  while True:
2      mutex.wait()
3      # critical section
4      mutex.signal()
```

- However, if A, B, and C are using a mutex in a loop, A and B could race around and around, starving C.

Semaphore starvation

- The weakest assumption about semaphores that makes it possible to avoid starvation is:
- **Property 3:** if there are threads waiting on a semaphore when a thread executes `signal`, then one of the waiting threads has to be woken.
- Prevents a thread from signalling a semaphore, racing around a loop and catching its own signal!

```
_____ Thread i _____  
1   while True:  
2       mutex.wait()  
3       # critical section  
4       mutex.signal()
```

- However, if A, B, and C are using a mutex in a loop, A and B could race around and around, starving C.
- A semaphore with Property 3 is called a **weak semaphore**.

Semaphore starvation

- **Property 4:** if a thread is waiting at a semaphore, then the number of threads that will be woken before it is bounded.

Semaphore starvation

- **Property 4:** if a thread is waiting at a semaphore, then the number of threads that will be woken before it is bounded.
- FIFO queues satisfy this property.

Semaphore starvation

- **Property 4:** if a thread is waiting at a semaphore, then the number of threads that will be woken before it is bounded.
- FIFO queues satisfy this property.
- A semaphore with Property 4 is called a **strong semaphore**.

Semaphore starvation

- **Property 4:** if a thread is waiting at a semaphore, then the number of threads that will be woken before it is bounded.
- FIFO queues satisfy this property.
- A semaphore with Property 4 is called a **strong semaphore**.
- Dijkstra (inventor of semaphores) conjectured in 1965 that it was impossible to solve the mutex problem without starvation with weak semaphores.

Semaphore starvation

- **Property 4:** if a thread is waiting at a semaphore, then the number of threads that will be woken before it is bounded.
- FIFO queues satisfy this property.
- A semaphore with Property 4 is called a **strong semaphore**.
- Dijkstra (inventor of semaphores) conjectured in 1965 that it was impossible to solve the mutex problem without starvation with weak semaphores.
- Morris showed you could do it in 1979.

Morris's algorithm

Initialization


```
1 room1 = room2 = 0
2 mutex = Semaphore(1)
3 t1 = Semaphore(1)
4 t2 = Semaphore(0)
```

```
1 mutex.wait()
2   room1 += 1
3 mutex.signal()
4 t1.wait()
5   room2 += 1
6   mutex.wait()
7   room1 -= 1
8   if room1 == 0:
9       mutex.signal()
10      t2.signal()
11 t2.wait()
12   room2 -= 1
13   # critical section
14   if room2 == 0:
15       t1.signal()
16   else:
17       t2.signal()
```

The Dining Philosophers

- Five philosophers are eating spaghetti.
- There are five forks.
- Eating spaghetti requires two forks.
- More than one philosopher can eat at a time.

```
1  _____ Philosopher i _____  
2  while True:  
3      think()  
4      get_forks()  
5      eat()  
      put_forks()
```



The Dining Philosophers: a Non-solution

Which fork?

```
1 def left(i) = return i
2 def right(i) = return (i+1)%5
```

Initialization

```
1 forks =
2     [Semaphore(1)
3     for i in range(5)]
```

Non-solution

```
1 def get_forks(i):
2     fork[right(i)].wait()
3     fork[left(i)].wait()
4
5 def put_forks(i):
6     fork[right(i)].signal()
7     fork[left(i)].signal()
```

- Why does this not work?

The Dining Philosophers: a Non-solution

Which fork?

```
1 def left(i) = return i
2 def right(i) = return (i+1)%5
```

Initialization

```
1 forks =
2     [Semaphore(1)
3     for i in range(5)]
```

Non-solution

```
1 def get_forks(i):
2     fork[right(i)].wait()
3     fork[left(i)].wait()
4
5 def put_forks(i):
6     fork[right(i)].signal()
7     fork[left(i)].signal()
```

- Why does this not work?
- Deadlock is possible.
How?

The Dining Philosophers: Solution #1

Initialization

```
1 footman = Semaphore(4)
```

Solution #1

```
1 def get_forks(i):  
2     footman.wait()  
3     fork[right(i)].wait()  
4     fork[left(i)].wait()  
5  
6 def put_forks(i):  
7     fork[right(i)].signal()  
8     fork[left(i)].signal()  
9     footman.signal()
```

- Only allow 4 philosophers at a time.
- Can you prove deadlock is impossible?

The Dining Philosophers: Solution #2

- Have at least one leftie and at least one rightie.
- How can you prove deadlock is impossible?

The Dining Philosophers: Tanenbaum's solution

Initialization

```
1 state = ['thinking'] * 5
2 sem = [Semaphore(0) for i in range(5)]
3 mutex = Semaphore(1)
```


```
1 def get_fork(i):
2     mutex.wait()
3     state[i] = 'hungry'
4     test(i)
5     mutex.signal()
6     sem[i].wait()
7
8 def put_fork(i):
9     mutex.wait()
10    state[i] = 'thinking'
11    test(right(i))
12    test(left(i))
13    mutex.signal()
```

```
1 def test(i):
2     if (
3         state[i] == 'hungry'
4         and
5         state(left(i)) != 'eating'
6         and
7         state(right(i)) != 'eating'
8     ) :
9         state[i] = 'eating'
10        sem[i].signal()
```

Tanenbaum's solution not starvation-free

- Can you think of a situation where a Tanenbaum philosopher can starve?


0



Tanenbaum's solution not starvation-free

- Can you think of a situation where a Tanenbaum philosopher can starve?

0



- Imagine trying to starve 0.
- Initially 1 and 3 are eating
- 3 swaps with 4
- 1 swaps with 2
- Now 2 and 4 are eating
- Repeat

Cigarette smokers problem

- To smoke you need: paper, tobacco, and a match.
- Four threads: an agent and three smokers.
- One smoker already has lots of paper.
- One smoker already has lots of tobacco.
- One smoker already has lots of matches.
- The agent repeatedly obtains two ingredients at random.
- The agent should wake up the smoker who needs those ingredients.

Cigarette smokers problem: agent code

- Interesting problem assumes we cannot change agent code.

Initialization

```
1 agentSem = Semaphore(1)
2 tobacco = Semaphore(0)
3 paper = Semaphore(0)
4 match = Semaphore(0)
```

Agent A

```
1 agentSem.wait()
2 tobacco.signal()
3 paper.signal()
```

Agent B

```
1 agentSem.wait()
2 paper.signal()
3 match.signal()
```

Agent C

```
1 agentSem.wait()
2 tobacco.signal()
3 match.signal()
```


Cigarette smokers problem: Non-solution

Agent A

```
1 agentSem.wait()
2 tobacco.signal()
3 paper.signal()
```

Agent B

```
1 agentSem.wait()
2 paper.signal()
3 match.signal()
```

Agent C

```
1 agentSem.wait()
2 tobacco.signal()
3 match.signal()
```

Smoker with matches

```
1 tobacco.wait()
2 paper.signal()
3 agentSem.signal()
```

Smoker with tobacco

```
1 paper.wait()
2 match.wait()
3 agentSem.signal()
```

Smoker with paper

```
1 tobacco.wait()
2 match.wait()
3 agentSem.signal()
```

- Why doesn't this work?

Smokers problem: Parnas solution

Initialization

```
1  isTobacco = isPaper =  
2      isMatch = False  
3  tobaccoSem = Semaphore(0)  
4  paperSem = Semaphore(0)  
5  matchSem = Semaphore(0)
```

Smoker with tobacco

```
1  tobaccoSem.wait()  
2  makeCigarette()  
3  agentSem.signal()  
4  smoke()
```

Pusher A

```
1  tobacco.wait()  
2  mutex.wait()  
3      if isPaper:  
4          isPaper = False  
5          matchSem.signal()  
6      elif isMatch:  
7          isMatch = False  
8          paperSem.signal()  
9      else:  
10         isTobacco = True  
11         mutex.signal()
```

Generalized smokers problem

Initialization

```
1 numTobacco = numPaper =  
2   numMatch = 1  
3 tobaccoSem = Semaphore(0)  
4 paperSem = Semaphore(0)  
5 matchSem = Semaphore(0)
```

Smoker with tobacco

```
1 tobaccoSem.wait()  
2 makeCigarette()  
3 agentSem.signal()  
4 smoke()
```

Pusher A

```
1 tobacco.wait()  
2 mutex.wait()  
3   if numPaper:  
4       isPaper -= 1  
5       matchSem.signal()  
6   elif numMatch:  
7       isMatch -= 1  
8       paperSem.signal()  
9   else:  
10       isTobacco += 1  
11   mutex.signal()
```

- Keep a **scoreboard**:
variables describing the state of the system.
- Each agent checks the scoreboard to see if it can proceed.